

Step 6 Completion Report

Adaptive Sub-Stepping — Before vs After Analysis

Neural Network Emulation of the Unrestricted Three-Body Problem

Aarush Gupta, S. Karthik Yadavalli · April 28, 2026

1. Executive Summary

Core finding: Adaptive distance-based sub-stepping eliminates catastrophic orbital instability (body ejection) and reduces worst-case trajectory error by **105×** (from 161.7 AU to 1.54 AU), while reducing the divergence rate from $\lambda = 0.251/\text{yr}$ to $\lambda = 0.156/\text{yr}$ (38% improvement). The cost of this stability is a 35% increase in per-step compute time (17.9 $\rightarrow$ 24.3 μs). This represents one of the most impactful engineering decisions in the entire SIMON development.

This report documents the before/after comparison for adaptive sub-stepping — a mechanism introduced between V0 and V1 of the paper. The comparison was not previously presented in any version of the paper (V0, V1, or V2). It quantifies the improvement using actual experimental data from two separate runs on the same hardware (NVIDIA RTX 5050).

2. Background — What Is Adaptive Sub-Stepping?

2.1 The Problem Without Adaptive Sub-Stepping

SIMON uses a fixed macro timestep dt to advance the simulation. During ordinary orbital evolution, a fixed dt works well — the gravitational force changes slowly and the leapfrog integrator remains accurate. However, in the three-body problem, bodies occasionally undergo close encounters where they pass within a fraction of an AU of each other. At these moments, gravitational acceleration changes dramatically within a single timestep.

The gravitational force between two bodies scales as $F \propto 1/r^2$. When $r \rightarrow 0$, $F \rightarrow \infty$. A fixed timestep that is appropriate for $r = 1$ AU is far too large when $r = 0.01$ AU — the force changes by a factor of 10,000 within the same interval. The integrator misses the true trajectory through the close encounter, introducing a large position error. In a chaotic system, this error amplifies exponentially, and the lightest body can acquire enough energy to escape the system entirely (ejection).

Without adaptive sub-stepping: Close encounters at certain dt values cause body ejection. The lightest body (0.005 $M_{\odot}$) reaches 161.7 AU — effectively lost from the system — when $dt = 0.01$ yr. This is physically non-real: the same system integrated at higher accuracy by ias15 shows no ejection.

2.2 The Mechanism — How Adaptive Sub-Stepping Works

At each macro timestep, SIMON computes the minimum pairwise separation $r_{\min}$ between all pairs of bodies. If $r_{\min}$ falls below a threshold (0.05 AU), the macro step is subdivided into n_{sub} smaller steps:

$$n_{\text{sub}} = \min(16, \text{ceil}(0.05 / r_{\text{min}}))$$

Each sub-step uses velocity-Verlet integration with duration dt/n_{sub} . As bodies approach each other, the number of sub-steps increases proportionally, so the effective resolution through close encounters scales automatically with the encounter severity. The cap of 16 sub-steps bounds the computational cost.

The mechanism activates on only 0.2–1.1% of all macro steps — the vast majority of steps require no sub-division. The overhead is therefore minimal: +35% in per-step time ($17.9 \rightarrow 24.3 \mu\text{s}$), averaged across all steps.

3. Experimental Setup

Both runs use identical conditions: same trained SIMON model (pair_correction_nn.pt), same initial conditions, same 5 dt values, same 100-year simulation horizon, same hardware (NVIDIA RTX 5050 CUDA 13.0). The only difference is whether the adaptive sub-stepping mechanism is active.

Parameter	Before Adaptive	After Adaptive
Adaptive threshold	0.0 AU (disabled)	0.05 AU (enabled)
Max sub-steps per macro step	N/A	16
dt values tested (yr)	[0.005, 0.01, 0.02, 0.04, 0.08]	[0.005, 0.01, 0.02, 0.04, 0.08]
Simulation horizon T	100 years	100 years
SIMON model	pair_correction_nn.pt (same)	pair_correction_nn.pt (same)
Hardware	NVIDIA RTX 5050 CUDA 13.0	NVIDIA RTX 5050 CUDA 13.0
Initial conditions m (M $\blacksquare$)	[1.0, 0.01, 0.005]	[1.0, 0.01, 0.005]

4. Results

4.1 Main Comparison Figure

Figure 1 shows the three key metrics side by side: final error vs timestep (left), divergence slope (centre), and per-step compute time (right).




Figure 1. Three-panel comparison of SIMON with and without adaptive sub-stepping. (a) Final RMS error vs ias15 at T=100yr across all dt values. The spike to 162 AU at dt=0.01 without adaptive stepping is eliminated entirely. (b) Divergence rate λ : 38% reduction with adaptive sub-stepping. (c) Per-step compute time: 35% overhead introduced by the mechanism.

4.2 Quantitative Summary

Metric	Before Adaptive	After Adaptive	Change
Bodies bounded at T=100yr?	NO — ejection to 162 AU	YES — all within 2.5 AU	Qualitative reversal
Divergence slope λ (1/yr)	0.2508	0.1562	−38%
Worst-case error (any dt)	161.7 AU (dt=0.01)	1.54 AU (dt=0.04)	105× lower
Error range across dt	[1.02, 161.7] AU	[0.847, 1.543] AU	Stable, tight range
Best speedup vs ias15	2.63× (dt=0.08)	2.09× (dt=0.08)	Slight reduction
Per-step time (dt=0.02)	17.9 μ s	24.3 μ s	+35% overhead
Adaptive activations	0% (disabled)	0.2–1.1% of steps	Minimal overhead

4.3 Divergence Rate Plots — Side by Side

Figures 2 and 3 show the SIMON-vs-ias15 RMS position error over the full 100-year simulation, before and after adaptive sub-stepping respectively. The visual difference is striking and tells the story clearly.




Figure 2. **Before** adaptive sub-stepping. The sharp discontinuous jump at ~year 30 is a body ejection signature. The error saturates at ~10 AU as the ejected body reaches large distances. $\lambda = 0.251/\text{yr}$.




Figure 3. **After** adaptive sub-stepping. Smooth exponential growth throughout. No discontinuities. The error plateaus at ~5 AU consistent with chaotic phase drift. $\lambda = 0.156/\text{yr}$.

4.4 Speed-Accuracy Frontier — Side by Side

Figures 4 and 5 show the speed-accuracy frontier for each configuration. The before plot reveals the catastrophic outlier: $dt=0.01$ produces a final error of 162 AU (body ejection), completely outside the expected chaotic saturation range. After adaptive sub-stepping, all dt values produce errors within [0.847, 1.543] AU — the tight cluster characteristic of chaotic saturation, not ejection.




Figure 4. **Before** adaptive sub-stepping. The $dt=0.01$ point shows 162 AU error — body ejection. The remaining points cluster near 1 AU but the $dt=0.01$ outlier makes the result unreliable.




Figure 5. **After** adaptive sub-stepping. All dt values produce errors within $[0.847, 1.543]$ AU. The frontier is clean and monotonic. No outliers. Best accuracy at $dt=0.01$ (0.847 AU).

5. Interpretation of Results

5.1 Why Does $dt=0.01$ Cause Ejection Without Adaptive Stepping?

This is the most important question for understanding the result. At $dt=0.01$ yr, each macro step spans ~ 3.65 days of simulation time. During a close encounter where two bodies pass within $r = 0.003$ AU of each other (which occurs in this initial condition), the gravitational acceleration changes by a factor of $\sim (0.01/0.003)^2 \approx 11$ within a single macro step. The leapfrog integrator assumes the force is roughly constant across each step. This assumption fails catastrophically during the close encounter, introducing a large spurious kick to the lightest body.

By contrast, $dt=0.005$ and $dt=0.02$ happen to avoid the worst part of the encounter trajectory and produce bounded results. $dt=0.04$ and $dt=0.08$ use the NN correction mechanism more aggressively and the gating system absorbs the error. $dt=0.01$ falls into the worst-case window for this particular initial condition.

Physical interpretation: Without adaptive sub-stepping, the simulation's accuracy during close encounters is entirely dependent on the relationship between the macro timestep and the encounter geometry. Certain dt values are accidentally safe; others trigger ejection. This is not a property of the physics — it is a numerical artifact. Adaptive sub-stepping eliminates this artifact by ensuring sufficient resolution at every close encounter regardless of the macro dt chosen.

5.2 Why Does the Divergence Rate Improve?

The reduction in divergence rate from $\lambda = 0.251/\text{yr}$ to $\lambda = 0.156/\text{yr}$ (38%) is a genuine physical improvement, not just elimination of the outlier. Even in configurations that do not eject ($dt = 0.005, 0.02, 0.04, 0.08$ before adaptive), the adaptive mechanism reduces the accumulation of integration error during close encounters. Each encounter is resolved more accurately, so the phase drift that drives chaotic divergence grows more slowly over the 100-year horizon.

The divergence rate fit uses the 10–50 year window, which captures the smooth exponential growth phase for both configurations. The before-adaptive curve shows a steeper slope in this window because even non-ejecting runs accumulate more phase error through imprecise close-encounter handling.

5.3 The Cost-Benefit Assessment

Adaptive sub-stepping adds 35% per-step overhead ($17.9 \rightarrow 24.3 \mu\text{s}$). This cost is incurred on every step, even though the mechanism only activates on 0.2–1.1% of steps. The overhead comes primarily from computing $r_{\min}$ at each macro step to check whether sub-division is needed.

- **What is gained:** 105× worst-case error reduction, 38% lower divergence rate, qualitative stability transition from ejection to bounded orbits.
- **What is lost:** 35% per-step speed. Best speedup vs ias15 drops from 2.63× to 2.09×.
- **Net verdict:** The trade is strongly favourable. A 35% speed cost for 105× accuracy improvement and elimination of non-physical ejections is an excellent engineering exchange. Without adaptive stepping, SIMON produces physically incorrect results at certain dt values — making those results unpublishable.

6. Integration into the V3 Paper

6.1 Why This Was Missing from V2

The V2 paper describes adaptive sub-stepping in Section 3.7 (methodology) and references the mechanism in Conclusion 2, but never presents the before/after comparison. A reader of V2 learns that adaptive sub-stepping exists and what it does mechanically, but has no quantitative evidence for why it was needed or how much it helps. This is a significant gap: the 105× improvement is the most dramatic single result in the project and is currently invisible.

6.2 New Paper Content for V3

Addition 1 — New Section 4.4: Adaptive Sub-Stepping Performance

Add a new Results subsection immediately before the Speed-Accuracy Frontier (current Section 4.3). This section should contain:

- The before/after table from this report (Table 1 above)
- Figure 1 (the 3-panel comparison figure) as the primary illustration
- One paragraph explaining the $dt=0.01$ ejection mechanism
- One sentence stating the net cost-benefit: 105× error reduction at 35% speed cost

Addition 2 — Expanded Conclusion 2

V2 Conclusion 2 currently reads: *"Adaptive sub-stepping eliminates non-physical behavior: The introduction of distance based dt adaptation during close encounters allows SIMON to preserve orbital dynamics for all three bodies across all tested timesteps."*

This should be expanded to include the specific numbers, which are currently absent:

"Adaptive sub-stepping eliminates non-physical ejections and reduces worst-case trajectory error by 105× (from 161.7 AU to 1.54 AU). Without this mechanism, SIMON produces body ejections at $dt = 0.01$ yr — a physically incorrect result caused by insufficient resolution during close encounters. With adaptive stepping active on just 0.2–1.1% of macro steps, all three bodies remain gravitationally bounded across all tested timesteps, at a cost of 35% additional per-step compute time (17.9 → 24.3 μ s)."

Addition 3 — Updated Section 3.7 Motivation

Section 3.7 currently describes the mechanism without explaining why it was needed. Add one sentence at the start: *"Without this mechanism, SIMON produces non-physical body ejections at certain timesteps, with worst-case position errors exceeding 161 AU for the 1, 0.01, 0.005 $M_{\odot}$ configuration at $T = 100$ yr."* This gives the mechanism its scientific motivation rather than presenting it as an arbitrary design choice.

6.3 Connection to the Broader Paper Narrative

In the context of the full V3 paper, the adaptive sub-stepping result connects naturally to the Phase 1 Vector NN finding. Both results make the same fundamental point: *physics-informed constraints are not optional refinements — they are necessary conditions for correctness.*

- **Phase 1 showed:** Removing analytic force direction (Vector NN) causes ejection through accumulated sub-degree directional errors.
- **Step 6 shows:** Removing adaptive close-encounter resolution causes ejection through insufficient integration accuracy.
- **Together they say:** SIMON's orbital stability is the product of two physics constraints working in combination — exact force direction and adaptive encounter resolution. Removing either one is sufficient to cause catastrophic failure.

Recommended framing for V3 abstract: *"We show that long-horizon orbital stability in chaotic three-body systems requires two complementary physical constraints: structurally enforced analytic force direction and adaptive timestep resolution during close encounters. Removing either constraint independently causes body ejection in 100-year simulations, demonstrating that physics-informed architectural choices are necessary, not merely beneficial, for gravitational neural network stability."*

End of Step 6 Report — Adaptive Sub-Stepping · SIMON V3 Research · April 28, 2026

Adaptive Sub-Stepping Evaluation — Correction Report

1. Purpose

This document reviews the before/after adaptive sub-stepping experiment and proposes targeted corrections to improve technical clarity, consistency, and reviewer robustness before final submission.

2. Summary of Findings

The adaptive sub-stepping experiment is well-designed and isolates a single variable: enabling vs disabling adaptive sub-stepping under identical simulation conditions. The results clearly demonstrate that adaptive sub-stepping eliminates timestep-induced instability and significantly improves accuracy.

Key validated outcomes:

- Worst-case error reduction: 161.7 AU $\rightarrow$ 1.543 AU ($\sim 105\times$ improvement)
- Divergence rate reduction: $\sim 0.251 \rightarrow \sim 0.156$ /yr ($\sim 38\%$ reduction)
- Stability improvement: Ejection $\rightarrow$ Fully bounded system
- Computational cost: +35% per-step overhead
- Activation frequency: only $\sim 0.2\text{--}1.1\%$ of steps

3. Identified Issues

Issue 1: Divergence rate inconsistency

Different parts of the project report slightly different post-adaptive divergence rates (~ 0.156 vs ~ 0.1655 /yr).

Issue 2: Interpretation wording

The report currently describes the improvement as a 'physical improvement', which is not technically precise.

4. Proposed Corrections

Correction 1: Clarify divergence rate reporting

Suggested wording:

The divergence rate after adaptive sub-stepping lies in the range 0.156–0.166 /yr depending on the fit window and summary run. The qualitative conclusion remains unchanged.

Correction 2: Replace interpretation language

Replace: genuine physical improvement

With: numerical stability improvement due to improved resolution of close encounters.

5. Strengthened Final Claim (Recommended)

Adaptive sub-stepping is the dominant stability mechanism in SIMON. It eliminates timestep-induced ejection during close encounters, reduces worst-case trajectory error by over two orders of magnitude, and lowers the divergence rate, while introducing only modest computational overhead and activating in a small fraction of simulation steps.

6. Final Recommendation

The experiment is strong and publication-ready after incorporating the above corrections. The results provide clear, quantitative evidence for the role of adaptive sub-stepping and significantly strengthen the paper's engineering contribution.