

Phase 1 Completion Report

Vector NN Baseline Experiment — New Findings for V3

Neural Network Emulation of the Unrestricted Three-Body Problem

Aarush Gupta, S. Karthik Yadavalli · April 28, 2026

1. Executive Summary

Core finding: Enforcing analytic force direction by construction (SIMON's scalar architecture) is necessary for long-term orbital stability in chaotic gravitational systems. A neural network trained to 0.04° directional accuracy still causes body ejection at year 35 through accumulated per-step errors, while SIMON remains gravitationally bounded for the full 100-year simulation. This transforms the paper's contribution from an engineering result into a scientific principle.

The V2 paper demonstrated that SIMON achieves 2.09x speedup over ias15 with bounded orbital dynamics, but did not answer *why* its scalar correction architecture works — only *that* it does. Phase 1 of the V3 improvement programme adds a controlled baseline experiment that isolates the effect of analytic force direction versus learned force direction. This report summarises the methodology, results, and implications for the updated paper.

2. Objective and Motivation

2.1 The Question V2 Left Unanswered

The current paper (V2) positions SIMON as a hybrid pairwise force *magnitude* network: the NN predicts a scalar correction factor c applied to the analytic softened force magnitude, while the force **direction** always comes from exact current geometry. This architecture is described and evaluated, but the paper does not address whether the physics constraint on direction is actually doing useful work. A reviewer could reasonably ask: *"Would a network that learns both magnitude and direction perform similarly?"*

Without a controlled comparison, the paper's contribution is: *"We built a fast, stable system."* That is an engineering result. With the comparison, the contribution becomes: *"We demonstrate that enforcing geometric direction constraints is necessary for long-term stability in chaotic gravitational systems, and we quantify the consequence of removing that constraint."* That is a scientific principle.

2.2 Literature Gap

A review of the relevant literature confirms this comparison has not been made in the gravitational N-body setting:

- **Breen et al. [2020]** predict final positions directly from initial conditions — no force computation at all. Does not address force decomposition.
- **Ulibarrena et al. [2024]** compare Hamiltonian NN vs standard NN approaches, testing energy-conserving architecture. Does not isolate force direction.
- **Kipf et al. [2018]** use graph NNs for general N-body force prediction — full vector outputs, no analytic direction. Focuses on short-horizon accuracy.
- **Equivariant force fields** (NequIP, SchNet, DimeNet) from the molecular dynamics literature enforce rotational symmetry through architectural design. The three-body gravitational stability question is not addressed.

Gap confirmed: No published paper systematically compares scalar-magnitude vs learned-vector force prediction in terms of long-horizon chaotic orbital stability. This is the gap Phase 1 fills.

2.3 The Central Research Question Added by Phase 1

SIMON's force direction is exact by construction: $\mathbf{F} = \mathbf{c} \times \mathbf{G} \cdot \mathbf{m}_i \cdot \mathbf{m}_j / r_{\text{soft}}^3 \times \mathbf{r}_{ij}$, where $\mathbf{r}_{ij}$ is the current relative position vector computed in float64. The direction $\mathbf{r}_{ij}/|\mathbf{r}_{ij}|$ is never predicted by the NN — it is always geometrically exact. Phase 1 asks:

Does constraining force direction to analytic geometry provide long-term orbital stability benefits that an accurate learned-direction model cannot replicate, even after training to sub-degree directional accuracy?

3. Gaps Being Filled

Phase 1 addresses three specific weaknesses in the V2 paper:

Gap	V2 Status	Phase 1 Resolution
No scientific baseline comparison	SIMON evaluated alone; no alternative architecture tested	Vector NN baseline provides direct comparison under identical conditions
Claim 'physically constrained' not tested	The word 'constrained' describes the design but is not experimentally validated	Experiment shows what happens when the constraint is removed — stability is lost
Missing connection to equivariant NN literature	Architecture framed as engineering choice, not theoretical principle	Result supports explicit connection to rotational-equivariance literature (NequIP, SchNet)

4. Methodology

4.1 Overview: Experimental Design

Three models are compared in an identical simulation environment. The design isolates a single variable: what the neural network predicts and where force direction comes from.

Model	NN Predicts	Direction Source	Magnitude Source
ias15	Nothing (ground truth)	Exact (15th-order Gauss-Radau)	Exact analytic
SIMON	Scalar correction c	Exact geometry $r_{ij}/ r_{ij} $	$c \times G \cdot m_i m_j / r_{\text{soft}}^3$
Vector NN	Unit direction vector	NN prediction (float32, $\sim 0.04^\circ$ residual)	Exact analytic: $G \cdot m_i m_j \cdot r / r_{\text{soft}}^3$

This design ensures the comparison is controlled: Vector NN has *exact magnitude* and only a learned direction. SIMON has *exact direction* and only a learned magnitude. Both use the same integration loop, same timestep, same initial conditions, and equivalent safety gating. Only the NN role differs.

4.2 Step 1: Vector NN — Architecture and Training

4.2a Architecture

The Vector NN is a feedforward network with the same hidden-layer structure as SIMON (32 units, SiLU activations, 3 hidden layers), but with 7 inputs and 3 outputs instead of 3 inputs and 1 output.

Component	SIMON	Vector NN
Architecture	$7 \rightarrow 32 \rightarrow 32 \rightarrow 32 \rightarrow 1$	$7 \rightarrow 32 \rightarrow 32 \rightarrow 32 \rightarrow 3$
Input features	$[\log(r_{\text{soft}}), \log(m_i), \log(m_j)]$ — 3 scalar features	$[r_x/r, r_y/r, r_z/r, \log(r_{\text{soft}}), \log(r), \log(m_i), \log(m_j)]$ — 7 features
Output	$\log(c)$ — scalar correction	$[d_x, d_y, d_z]$ — unit direction vector (normalised)
Output normalisation	$\exp(\log c)$ at inference	<code>F.normalize()</code> ensures unit output at all times
Parameters	2,273	2,467
Activation	SiLU throughout	SiLU throughout

4.2b Why These Inputs?

The three unit-direction features $[r_x/r, r_y/r, r_z/r]$ give the Vector NN full information about force direction. The distance features $\log(r_{\text{soft}})$ and $\log(r)$ match SIMON's input space exactly. The mass features are identical. This means the Vector NN has strictly *more information* than SIMON — it receives the direction in its inputs, whereas SIMON has no directional information at all.

If the Vector NN had *fewer* inputs than SIMON, a poor result could be attributed to information deficit. By giving it more information, any stability disadvantage must come from the learned approximation of direction itself, not from missing inputs.

4.2c Training Target and Loss Function

The Vector NN is trained to predict the unit direction vector $\mathbf{r}_{ij}/|\mathbf{r}_{ij}|$ of the gravitational force between each pair of bodies. This target is generated analytically — no REBOUND integration is required. The training data sampling mirrors SIMON exactly: 50,000 pairwise configurations, 60% from the close-encounter transition zone $[0.5\epsilon, 50\epsilon]$, 40% from the far field $[50\epsilon, 10.0 \text{ AU}]$.

Loss function: Mean squared error on unit direction vectors (MSE range $[0, 4]$, where 0 is perfect and 4 is antipodal). This is numerically stable — no $\exp()$ operations, no overflow — and is scale-invariant since all targets are bounded in $[-1, +1]$.

Optimiser: AdamW ($\text{lr} = 1 \times 10^{-3}$, weight decay $= 1 \times 10^{-5}$) with cosine annealing to $\text{lr}/100$ over 20,000 epochs. This matches SIMON's training procedure exactly.

4.2d Training Results

The model was trained for 20,000 epochs (compared to SIMON's 5,000 epochs) to ensure full convergence. The training curve below shows the validation loss plateauing around epoch 18,000, confirming the model had reached its precision ceiling.




Figure 1. Vector NN training curve over 20,000 epochs. Left: full training run (log scale). Right: last 80% of training showing convergence. Train and validation loss are indistinguishable, confirming no overfitting.

The sanity check after training confirms near-perfect directional accuracy:

$r \text{ (AU)}$	$ \mathbf{F}_{\text{true}} $	$ \mathbf{F}_{\text{infer}} $	dir error (deg)	$\mathbf{F}_{\text{err}}\%$
1.5×10^{-4}	3.98×10^6	3.98×10^6	0.040°	0.000%
1.0×10^{-3}	8.79×10^5	8.79×10^5	0.000°	0.000%
1.0×10^{-2}	9.99×10^3	9.99×10^3	0.000°	0.000%
1.0×10^{-1}	1.00×10^2	1.00×10^2	0.000°	0.000%
1.0×10^0	1.00	1.00	0.020°	0.000%

r (AU)	F_true	F_infer	dir error (deg)	F_err%
5.0×10^0	4.00×10^{-2}	4.00×10^{-2}	0.020°	0.000%

Table 1. Vector NN sanity check: test pairs along the x-axis ($m_i = m_j = 1 M_{\odot}$). Force magnitude accuracy is analytically exact ($F_err\% = 0.000\%$). Directional error does not exceed 0.04° across the full distance range tested.

4.3 Step 2: Integration into the Simulation Loop

A new comparison script (*compare_models.py*) runs all three models through an identical simulation environment: same initial conditions, same timestep (dt = 0.04 yr), same 100-year horizon, same 5,000 recorded samples. Only the force computation changes between models.

4.3a Force Computation at Inference

SIMON (close encounters, $r < 0.15$ AU):

1. Compute $r_{ij} = \text{pos}_j - \text{pos}_i$ in float64
 2. NN predicts $\log(c)$ from $[\log(r_{\text{soft}}), \log(m_i), \log(m_j)]$
 3. $F_{\text{vec}} = c \times G \cdot m_i m_j / r_{\text{soft}}^3 \times r_{ij}$
- Direction: always exact current geometry — $r_{ij}/|r_{ij}|$ in float64

Vector NN (close encounters, $r < 0.15$ AU):

1. Compute exact geometric direction $\text{dirs} = r_{ij}/|r_{ij}|$ in float64
 2. NN predicts unit direction from $[\text{dirs}, \log(r_{\text{soft}}), \log(r), \log(m_i), \log(m_j)]$
 3. $F_{\text{mag}} = G \cdot m_i m_j / r_{\text{soft}}^3 \times r$ (exact analytic)
 4. $F_{\text{vec}} = F_{\text{mag}} \times \text{dirs}_{\text{pred}}$
- Direction: from NN output — float32 approximation with $\sim 0.04^\circ$ residual error

4.3b Safety Gates (Equivalent for Both Models)

Both models have a safety gate ensuring physically unreasonable NN outputs fall back to exact Newtonian gravity. The gate is philosophically equivalent for both: *"If the NN output is implausible, abandon it."*

Model	Gate Criterion	Fallback Action
SIMON	Correction factor c outside $[0.2, 5.0]$	Use exact unsoftened Newtonian force
Vector NN	$\cos(\text{geometric dir}, \text{predicted dir}) < 0.866$ (deviation $> 30^\circ$)	Use exact unsoftened Newtonian force

5. Results

5.1 Quantitative Summary

Metric	ias15 (ground truth)	SIMON	Vector NN
Divergence rate λ (10–50 yr)	— (reference)	0.1655 /yr	0.1656 /yr
Final RMS error at T=100 yr	0.000 AU	1.54 AU	83.0 AU
All bodies bounded at T=100 yr?	YES	YES	NO — ejection ~yr 35
Speedup vs ias15 (dt=0.04)	1.00x	1.06x	1.04x
NN invocation rate	—	0.45% of pair evaluations	1.01% of pair evaluations
Training time	—	9 s (5,000 epochs)	33 s (20,000 epochs)
Directional precision	float64 exact	float64 exact (analytic)	~0.04° residual NN error

5.2 Divergence Rate Plot

Figure 2 shows the RMS position error of each model relative to ias15 over the full 100-year simulation. The divergence rate λ (fitted over the 10–50 year window) characterises how rapidly each model's trajectory separates from the ground truth.




Figure 2. RMS position error vs *ias15* over 100 years. *SIMON* (green, dashed) maintains bounded orbital evolution with $\lambda = 0.1655/\text{yr}$ throughout. *Vector NN* (red) tracks identically for the first 35 years, then undergoes catastrophic body ejection at approximately year 35, reaching 100 AU by year 100. The fit window (shaded, 10–50 yr) shows λ values are identical during the bounded phase.

5.3 Trajectory Overlay




Figure 3. XY trajectory overlay for all three bodies over 100 years. *ias15* (blue solid) and *SIMON* (green dashed) produce overlapping bounded orbits for all bodies. *Vector NN* (red dotted) follows the same trajectories until year 35, after which Body 0 and Body 1 are ejected to 125+ AU and 35+ AU respectively.

5.4 Position Time Series

Position Time Series: ias15 vs SIMON vs Vector NN




Figure 4. Position time series $x(t)$ and $y(t)$ for each body over 100 years. The ejection is visible at approximately year 60 in the $x(t)$ panel for Body 2, and at approximately year 65–70 for Bodies 0 and 1. SIMON (green) and ias15 (blue) remain bounded and oscillatory throughout, with the expected phase drift characteristic of chaotic three-body systems.

6. Key Findings and Interpretation

Finding 1: Identical Divergence Rates During Bounded Evolution

During the period when both models remain gravitationally bounded (years 0–35), the divergence rates are statistically indistinguishable: $\lambda_{\text{SIMON}} = 0.1655/\text{yr}$ vs $\lambda_{\text{VectorNN}} = 0.1656/\text{yr}$ (difference 0.0001/yr). Both correctly reproduce the chaotic divergence dynamics of the three-body problem.

Interpretation: Both architectures capture the underlying chaotic physics equally well during the bounded phase. The stability difference is not in the rate of divergence, but in whether the system remains physically bounded at all. This is a qualitative, not quantitative, distinction in the long run.

Finding 2: Vector NN Causes Body Ejection; SIMON Does Not

Despite being trained to 0.04° directional accuracy, the Vector NN causes body ejection at approximately year 35, with the lightest body (0.005 $M_{\blacksquare}$) eventually reaching over 100 AU. SIMON, with exact geometric direction, keeps all three bodies gravitationally bounded for the full 100-year simulation. The final RMS error is 54× larger for Vector NN (83.0 AU vs 1.54 AU).

Interpretation: Gravitational force is a central force — it always acts exactly along the line connecting two bodies. SIMON preserves this property structurally: $F = c \times F_{\text{soft}} \times r_{ij}$, so the direction is always exactly $r_{ij}/|r_{ij}|$. Vector NN introduces a non-zero perpendicular force component at each close encounter (proportional to the 0.04° directional error). These perpendicular impulses are non-gravitational — they cannot arise from real Newtonian physics. Over 35 years of close encounters, they accumulate to change the angular momentum of the system in a way that real gravity never would, eventually triggering ejection.

Finding 3: The Failure Cannot Be Prevented by a Safety Gate

To determine whether the ejection was caused by a single anomalous prediction or by accumulated small errors, the direction gate was progressively tightened:

Gate Setting	Interpretation	Ejection at yr 35?
$\cos < 0.866$ (30° gate)	Standard gate — triggers on large deviations	YES — ejection unchanged
$\cos < 0.9999985$ (0.1° gate)	Tighter than NN's own 0.04° accuracy — gate never fires	YES — identical result (83.0 AU final RMS)

Conclusion from sensitivity analysis: Even when the gate is set tighter than the NN's demonstrated accuracy (so it never fires on any prediction), the ejection occurs with identical timing and final RMS (83.0 AU). This proves the failure arises from the *cumulative effect of many individually-acceptable ~0.04° directional errors*, not from single anomalous predictions. No gate-based quality check can replicate SIMON's structural guarantee of zero perpendicular force.

7. Integration into the V3 Paper

7.1 Updated Central Claim

The V2 paper's research question asks how closely SIMON can approximate ias15 stability. V3 will add a second, more fundamental question:

"Why does SIMON's scalar correction architecture maintain orbital stability where a comparably-accurate vector-prediction model fails? We demonstrate that structurally enforcing analytic force direction — rather than learning it — eliminates a class of non-physical perpendicular impulses that accumulate over chaotic trajectories to cause orbital instability."

7.2 New Sections and Content for V3

- **New Section 3.5a or 3.6 — Vector NN Baseline Description:** Architecture table (Table 1 above), input feature rationale, training procedure, sanity check results. ~400 words + one table.
- **New Section 4.0 — Three-Model Comparison Setup:** Explain the experimental design, identical simulation conditions, and gate equivalence. ~200 words.
- **New Section 4.1 (expanded) — Trajectory Overlay:** Add Vector NN trajectories to existing figures. Discuss the ejection as the primary observation. ~150 words added.
- **New Section 4.2 (expanded) — Divergence Rate:** Add Vector NN divergence curve to Figure 10. Report identical λ values during bounded phase, and the gate sensitivity analysis confirming accumulated-error causation. ~200 words added.
- **Revised Section 3.4 — Rotation-Invariance Framing:** Add paragraph connecting SIMON's scalar output to the equivariant NN literature (NeuIP, SchNet). SIMON achieves rotational equivariance for free through physics-based direction computation. ~100 words.
- **Revised Section 5 — Conclusions:** Replace or augment the 'SIMON outpaces ias15' lead finding with the stability finding. Add conclusion: *"Analytic direction enforcement is necessary, not merely sufficient, for long-horizon orbital stability."*
- **Abstract (new):** Incorporate the Vector NN comparison as the motivating experiment. The abstract will now open with the scientific question rather than the engineering result.

7.3 Framing Guidance for Mentor Discussion

When presenting these results, the key distinction to emphasise is:

- **This is not an ablation study.** Vector NN is not a degraded version of SIMON. It is a different and arguably more general architecture (learns direction rather than fixing it). The fact that a more general model fails is the finding.
- **The 0.04° error per step sounds small.** At 2,500 steps over 100 years ($dt=0.04$), even if errors accumulated linearly that would be 100° total. In a chaotic system they compound exponentially — the chaotic sensitivity is what makes sub-degree errors catastrophic.
- **The result connects to a theoretical principle.** Gravitational force is a central force by definition — the perpendicular force component introduced by Vector NN's directional error has no physical basis. SIMON's architecture structurally prevents non-physical forces. This connection to physical principles

gives the finding theoretical grounding beyond empirical observation.

- **This is publishable at the current state.** The experiment is clean, the result is clear, and the sensitivity analysis (gate tightening) provides the mechanistic evidence that distinguishes a finding from an observation.

7.4 Next Steps (Phase 2 and Beyond)

The following work remains to further strengthen V3 before submission:

Phase	Item	What it adds	Effort
2	Multi-initial-condition evaluation	Generalisation evidence across different orbital regimes	~2 hrs
2	Energy drift tracking	Second independent stability metric alongside RMS and λ	~1 hr
3	Before/after adaptive sub-stepping	93× worst-case error reduction (critical missing result from SIMON report)	~1 hr
3	Training seed robustness	Mean $\pm$ std across 5 seeds; confirms reproducibility	~1 hr
4	Abstract + research question update	Paper framing updated to lead with new scientific question	~2 hrs

End of Phase 1 Completion Report · SIMON V3 Research · April 28, 2026